CLEARFORGE WEEKLY AI LEARNING BRIEF

From AI demos to usable systems

The practical upgrades, risks and experiments that mattered from 13–19 July 2026.

Human-led. AI-empowered.

Turning human input into clear, usable systems.

Mixed beginner + intermediate edition • Last verified: 20 July 2026, 06:45 BST

START HERE

How to use this learning brief

The one-hour route

Read the executive summary, choose one topic marked Use now or Test carefully, complete its guided exercise, and record what happened. You do not need to try every tool. The value is in making one measured change to a real workflow.

The evidence labels

Confirmed release means a feature is documented as available. **Staged rollout** means access is expanding but not universal. **Research preview** means the vendor is inviting early testing, not promising production reliability. **Vendor claim** means performance or impact has not been independently reproduced here. **Third-party reporting** is clearly identified.

The Clearforge verdicts

Use now: mature enough for a bounded task. **Test carefully**: useful, but access, quality or control needs checking. **Watch**: important direction, limited immediate action. **Skip for now**: cost, access or uncertainty outweighs likely value for most small teams.

- Keep a human approval point before publishing, paying, deleting, sending or exposing private data.
- Measure the whole task, including correction time—not just how fast the first draft appears.
- Recheck the linked primary source before buying or deploying; plans and rollouts can change.

THE WEEK IN ONE VIEW

Executive summary

This week did not produce one single breakthrough that everyone must adopt. It produced something more useful: evidence that AI products are becoming connected systems that research, create, schedule, run code, scan work and measure outcomes. That makes them more capable—and raises the cost of weak permissions, vague success measures and unreviewed automation.

Development	Status	Clearforge call
Canva AI 2.0	Research preview; broader rollout promised	Test carefully
Inkling open weights	Released; developer-focused	Watch / test if technical
GitHub AI security detections	Public preview; paid requirements	Test carefully
Gemini Notebook	Rename + staged code execution rollout	Use now for eligible users
GPT-Red	Published safety research; internal system	Watch and adopt the lesson
AI workflow scorecard	Vendor framework, immediately usable	Use now
EU Android/search measures	Binding regulatory measures	Watch
ASML capacity expansion	Confirmed company outlook	Watch

Central lesson: the best upgrade this week is not a new subscription. It is a better operating rule: define “done,” test on real inputs, keep review boundaries, and compare cost per successful outcome.

PLAIN-ENGLISH CONCEPT

The central development: AI becomes a workflow layer

What changed

Canva described a creative system that can connect to work tools, research the web, schedule tasks and build editable designs. Google added code execution to notebooks. GitHub placed AI detections inside pull requests. OpenAI published a scorecard centred on completed work rather than token volume. These are different products, but they point in the same direction: AI is moving from “answer this” toward “help move this task through a system.”

Why this matters

A chatbot can be checked one answer at a time. A workflow layer can touch files, calendars, customer messages, designs, code and publishing channels. That can remove friction, but one bad assumption can travel farther. Small businesses therefore need lightweight controls before they need more automation.

A simple workflow control pattern

Input: identify the source material. **Bound:** state what the AI may and may not do. **Generate:** produce a draft or analysis. **Verify:** check facts, permissions, brand and sensitive data. **Approve:** a named human decides whether the work advances. **Log:** record the result, cost and rework.

- Good first automation: assemble a draft briefing from approved sources.
- Poor first automation: publish claims, spend money or message customers without review.
- Success measure: useful output accepted with minimal correction—not raw output volume.

CLEARFORGE VERDICT — USE NOW: apply the control pattern to one existing AI-assisted task before adding another tool.

CREATOR TOOLS • 16 JULY

Canva AI 2.0: from design tool to creative operating system

What happened

Canva announced Canva AI 2.0 as a **research preview**. Canva says it can create layered editable work through conversation and adds six workflow areas: connectors, scheduling, web research, brand intelligence, Sheets AI and Canva Code 2.0. The company says general availability will roll out over the coming weeks; that is not the same as universal access today.

What it is

Instead of generating one flattened image, Canva says the system understands design objects, layout and hierarchy, allowing individual elements to remain editable. Connectors are intended to draw context from services such as Gmail, Google Drive, Slack, Notion and HubSpot. Scheduling is intended to run recurring creative tasks in the background.

Access, plans and limits

The announcement labels the product a research preview and does not provide a complete plan-by-plan price table for every capability. Early access and general rollout are separate. Treat any access you see as account-specific until Canva's live plan page confirms otherwise.

Previous position

Canva already offered conversational design and AI editing. The verifiable change is the proposed coordination layer: persistent memory, connectors, scheduled tasks, research and more structured outputs in one loop.

CLEARFORGE VERDICT — TEST CAREFULLY: high practical potential, but preview access and broad vendor claims require controlled testing.

GUIDED EXERCISE

Canva AI 2.0: practical test

Creator example

A solo creator has a launch brief, customer questions and brand guidelines. A useful test is not “make me a campaign.” It is: produce one editable launch graphic, one email outline and one seven-day content plan from the same approved brief, then check brand consistency and factual accuracy.

Example prompt
Using only the attached approved launch brief and brand guide, create an editable campaign starter set: one square announcement, one email outline and a seven-day content plan. Do not invent prices, testimonials, deadlines or product features. Mark any missing fact as NEEDS HUMAN INPUT. Keep every design element editable.

Human-review checks

Confirm the connector did not pull unrelated private material. Verify all claims against the source brief. Check accessibility, contrast, cropping and mobile readability. Confirm scheduled outputs remain drafts until a human approves them.

Limitations and risks

Connector permissions expand the data surface. Persistent memory can preserve outdated brand rules. Web research can introduce weak sources. Scheduling can repeat errors at scale. “On brand” remains a vendor claim until your own reviewers agree.

- Try one output format first.
- Record time to accepted result, including edits.
- Disconnect unused integrations after the test.
- Do not schedule publishing during the preview test.

MODELS • 15 JULY

Inkling: a major open-weight multimodal release

What happened

Thinking Machines Lab released Inkling with full weights available. The company describes a mixture-of-experts model with 975 billion total parameters, 41 billion active parameters and up to a one-million-token context window. It says the model was trained on text, images, audio and video and is available for fine-tuning through Tinker. These technical and benchmark statements are **vendor claims**.

What “open weights” means

The learned numerical parameters are available for others to inspect, host or customise under the stated licence and delivery terms. Open weights do not automatically mean open training data, low hardware cost, no restrictions or safer output.

Who can use it

The immediate audience is developers and organisations with model-deployment skills. The vendor also offers a Tinker playground and deployment partners. For a typical creator, the practical value may arrive through products built on Inkling rather than self-hosting it.

What changed

The release adds another large multimodal customisation option. The company itself says Inkling is not the strongest overall model; its pitch is a balanced base for customisation rather than benchmark supremacy.

CLEARFORGE VERDICT — WATCH: important for model choice and vendor independence; most non-technical creators should not self-host yet.

DECISION EXERCISE

Inkling: when customisation is worth it

Small-business example

A specialist support company wants an assistant that follows a narrow taxonomy and works with confidential internal material. Before fine-tuning, it should test whether retrieval over approved documents already solves the problem. Fine-tuning is justified only if repeatable behaviour remains poor after simpler controls.

Decision prompt
List the behaviour we need the model to perform repeatedly, the evidence that prompting and retrieval are insufficient, the acceptable error rate, the hosting boundary, the data we are permitted to use, the evaluation set, the rollback method and the named human owner. If any field is unknown, recommend a smaller test instead of fine-tuning.

Human-review checks

Read the licence and model card. Test factuality, refusal behaviour and data leakage on your own evaluation set. Price hosting, monitoring, updates and incident response—not just inference. Keep the original model and rollback path.

Risks

Large context can encourage users to upload too much data. Open weights move more safety and operations responsibility to the deployer. Vendor benchmarks may not match your task. Fine-tuning can reinforce bad examples.

- Use a 30–50 case evaluation set before and after customisation.
- Separate public, internal and restricted data.
- Log model version, prompt, retrieval set and reviewer outcome.

AUTOMATION + SAFETY • 14 JULY

GitHub brings AI security findings into pull requests

What happened

GitHub launched AI-powered security detections in **public preview**. Findings appear inside pull requests, run when a pull request is opened or updated, and are labelled as AI-generated. GitHub states they are informational and do not block merges.

Access and billing

The preview requires GitHub Code Security, CodeQL default setup, enterprise permission and a GitHub Copilot licence. Runs consume organisational AI credits. This is not a free security scanner for every repository.

What changed

CodeQL covers defined languages and query packs. GitHub says the AI layer broadens coverage to languages and frameworks beyond current CodeQL support. It depends on CodeQL setup but does not replace the underlying analysis or human security review.

Why it matters

The useful shift is placement. Advice delivered at the pull-request review point is more actionable than a separate report discovered after release. Small teams can apply the same lesson even without this product: put checks where decisions are made.

CLEARFORGE VERDICT — TEST CAREFULLY: use as another signal, never as proof that code is safe.

AI security review: a safe trial

Developer test

Choose a non-critical repository with known test cases. Enable the preview only if your plan and policy allow it. Open a pull request containing one benign known weakness in a test branch. Record whether the system finds it, how specific the explanation is and how many false positives appear.

Review prompt for a human
For each AI-labelled finding, identify the affected line, attack path, evidence, severity, safe reproduction method, proposed fix and test that proves the fix. If evidence is missing, mark the finding UNCONFIRMED rather than closing or merging automatically.

Human-review checks

Do not paste secrets into comments. Confirm licence and credit consumption. Require a developer to reproduce important findings. Keep dependency scanning, secret scanning, tests and code review. Track false positives so teams do not learn to ignore alerts.

Creator translation

Even if you never write code, the pattern is useful: attach AI review to the approval point. For a newsletter, that means citation and disclosure checks before send. For a product page, it means claims and price checks before publish.

RESEARCH TOOLS • 16 JULY

NotebookLM becomes Gemini Notebook—and gains code execution

What happened

Google renamed NotebookLM to Gemini Notebook. The standalone research product remains, but Google says notebooks will connect more closely with the Gemini app and eventually Search. A secure cloud computer can write and execute code for analysis grounded in notebook sources.

Access and rollout

Code execution is available now to Google AI Ultra users and Workspace business customers with AI Ultra Access or AI Expanded Access. Google says it will roll out to all Pro users on the web over the coming weeks. That is a staged rollout, not universal availability.

Previous position

NotebookLM was already a source-grounded research and learning tool. The significant addition is an execution environment: the notebook can analyse data with code, not only summarise and synthesise text. Cross-app syncing also moves it deeper into Google's ecosystem.

Why it matters

A notebook can now move from "what do these sources say?" toward "calculate, compare and visualise what these files contain." That is useful for survey results, product inventories and content audits—but executed code introduces new failure modes.

CLEARFORGE VERDICT — USE NOW: eligible users can test it on non-sensitive data with calculation checks; others should watch the rollout.

TRY IT YOURSELF

Gemini Notebook: a grounded analysis exercise

Small-business example

Upload an approved CSV of anonymous support categories and a short document defining each category. Ask the notebook to count requests by category, show uncertainty, and produce a short improvement brief grounded only in those sources.

Example prompt
Using only the attached files, calculate the number and percentage of requests in each approved category. Show the code or calculation method. Flag rows that do not match the taxonomy instead of guessing. Produce a table, a simple chart and three operational observations. Do not infer customer demographics or intent.

Human-review checks

Recalculate a sample manually. Inspect category mapping and missing values. Confirm the notebook did not treat narrative text as numeric fact. Avoid personal or regulated data unless your organisation has approved the service and configuration.

Limitations

Google controls the execution environment and rollout. Generated code can be wrong even when it runs. Source grounding reduces—but does not remove—misinterpretation. Cross-app syncing may affect where information appears, so review sharing and account settings.

A better AI scorecard: cost per successful task

What happened

OpenAI published a vendor framework arguing that useful work, cost per successful task and dependability matter more than token price alone. It recommends defining “done” in the system where work happens and includes human review, retries and rework in the full cost.

Why it matters

This is not a neutral industry standard; it is OpenAI’s framing and also supports its tiered model family. But the measurement method is practical and vendor-independent. A cheap model that needs three retries may cost more than a stronger model that succeeds once.

Formula
Full task cost ÷ number of outputs that meet the quality bar = cost per successful task.

Guided exercise

Choose one repeated task and run ten normal examples. Record model/API cost, staff minutes, retries, corrections and accepted outcomes. Compare against the existing non-AI method. Do not count a draft as successful unless it meets the real acceptance standard.

- Define success before running the test.
- Include failed attempts and review time.
- Use representative inputs, not showcase examples.
- Stop if quality or privacy falls below the existing process.

CLEARFORGE VERDICT — USE NOW: the framework is simple, cheap and useful regardless of vendor.

RESEARCH + SAFETY • 15 JULY

GPT-Red: automated red-teaming enters the training loop

What happened

OpenAI published research on GPT-Red, an internal automated red-teaming model used to find vulnerabilities and generate adversarial examples. OpenAI says it used GPT-Red to adversarially train GPT-5.6 against prompt injection. GPT-Red itself is not presented as a public customer product.

Plain-English explanation

Prompt injection is a malicious or misleading instruction hidden in data the AI reads—such as a webpage, email, file or tool response. It attempts to make the model ignore its intended rules, reveal data or take an unsafe action. Automated red-teaming creates many attack variations so weaknesses can be found before deployment.

What is confirmed versus claimed

Confirmed: OpenAI published the method and says it is part of its internal safety work. Vendor claim: the reported robustness gains and evaluation results. Unknown here: how the system performs across every real-world integration or against future attacks.

Practical action

Treat external content as untrusted. Separate read-only research from action tools. Require approval before sending, posting, purchasing, deleting or uploading data. Test your own workflows with harmless prompt-injection simulations.

CLEARFORGE VERDICT — WATCH: the research is important; the immediate value is the control lesson, not access to GPT-Red.

CONSUMER SAFETY • 16 JULY

Age prediction and teen safeguards become product infrastructure

What happened

OpenAI said it is strengthening teen protections and rolling out age prediction on consumer plans. The system estimates whether an account likely belongs to someone under 18 using account and behavioural signals, then applies additional protections. Users placed incorrectly can verify age through a third-party identity service.

Why it matters

This is a product-design shift: family and age controls are becoming part of everyday AI access rather than a separate parental guide. It may reduce exposure to sensitive content, while also raising questions about classification accuracy, privacy and appeals.

For creators and small businesses

If your product, community or educational content may reach minors, document the intended age range, moderation boundaries, escalation route and data minimisation. Do not treat a platform’s age estimate as proof of identity or as permission to collect more data.

Guided check

Review one audience-facing AI feature. Can a young user understand what it does? Are sensitive topics handled conservatively? Is there a visible route to a human? Can an incorrect restriction be appealed? Are parents or guardians given appropriate—but not excessive—control?

CLEARFORGE VERDICT — WATCH: important safety direction; most users should review settings rather than attempt to build age prediction themselves.

REGULATION • 16 JULY

EU measures push Android and search toward AI competition

What happened

The European Commission issued two sets of binding specification measures to Google under the Digital Markets Act. One aims to give competing AI services equal access to Android features used by Google's own services. The second sets conditions for sharing anonymised search data with eligible third-party search engines, including AI chatbots with search functions.

What it could change

The Commission says preferred AI assistants should be able to activate by voice and perform tasks in apps, subject to security and privacy safeguards. Search-data access could help competing services improve. Google has raised privacy and security objections, as reported by Reuters.

Who is affected

Immediate implementation work falls on Google and eligible providers. Creators and small businesses may eventually see more assistant choice and different search/discovery routes in the EU. This is not a promise that every assistant will gain every feature immediately.

Confirmed future milestones

Reuters reported implementation of the search-data measure from **January 2027** and user-facing Android benefits from **July 2027**. Treat those as confirmed reported milestones but recheck the Commission and Google before the date because implementation can change.

CLEARFORGE VERDICT — WATCH: strategically important, but there is little for a normal user to configure today.

The infrastructure signal: AI demand is still expanding

What happened

ASML reported second-quarter 2026 results and raised its 2026 sales outlook to €43–45 billion. The company linked demand to ongoing AI investment and said it plans to add 30% to 2026 low-NA EUV and DUV immersion capacity for 2027, while investigating another 30% increase for 2028.

What ASML does

ASML builds lithography systems used in advanced semiconductor manufacturing. It does not sell a chatbot. Its capacity plans are a useful upstream indicator of how chipmakers expect demand to develop.

Why creators should care

Infrastructure expansion does not guarantee lower retail AI prices. It does suggest vendors expect sustained demand. For small teams, AI should be budgeted as an operating cost with usage limits, fallback tools and a clear link to outcomes.

Cost-control exercise

List every AI subscription and API used last month. Map each to one workflow and one measurable outcome. Remove duplicates, set spending alerts, and keep a fallback for the one workflow that would hurt most if a service failed.

CLEARFORGE VERDICT — WATCH: use the signal for budgeting and vendor-risk planning, not investment advice.

Build a small, human-led AI stack

A sensible starter stack

Research: one source-grounded notebook. **Creation:** one design/content tool. **Automation:** one bounded handoff layer. **Review:** a human checklist at the decision point. **Evidence:** a simple log of source, output, cost and approval. Adding more tools before these roles are clear usually increases duplication.

Layer	Question	Minimum control
Source	What may the system read?	Approved folders and links
Model	What task is it chosen for?	Named use case and fallback
Action	What may it change?	Draft-only by default
Review	Who decides “done”?	Named human owner
Evidence	Can we reproduce the result?	Date, version, sources, outcome
Cost	Is it worth repeating?	Cost per accepted task

Rule of thumb: automate preparation before judgement. Let AI collect, sort, draft and compare. Keep a person responsible for claims, permissions, sensitive data and public action.

Terminology without the jargon

Agentic — Able to plan or take multiple tool-assisted steps toward a goal.

Context window — The amount of information a model can consider in one interaction.

Connector — A permissioned link between an AI system and another service or data source.

Fine-tuning — Further training a base model on examples to change behaviour or capability.

Foundation model — A broadly trained model that can support many downstream tasks.

Mixture of experts — An architecture that routes a task through selected parts of a larger model.

Multimodal — Able to work across more than one data type, such as text, image or audio.

Open weights — Model parameters are available; this does not necessarily mean fully open source.

Prompt injection — Instructions hidden in external content that try to redirect an AI system.

Red-teaming — Structured attempts to find failures, misuse paths or vulnerabilities.

Retrieval — Supplying relevant source material to a model at request time.

Staged rollout — Access expands gradually by account, plan, platform or region.

Research preview — Early access intended for learning and feedback, not a reliability guarantee.

Cost per successful task — Total cost—including review and retries—divided by accepted outcomes.

CHECK YOUR UNDERSTANDING

Weekly knowledge recap

1. Why is Canva AI 2.0 not a universal “use now” recommendation?

It is a research preview with access and reliability still rolling out.

2. Does open weights mean free, fully open and safe?

No. Licence, hosting cost, training data and deployment responsibility remain separate questions.

3. Why should AI security findings not block or approve a merge automatically?

They are informational signals that require evidence and human reproduction.

4. What changed in Gemini Notebook?

Renaming, deeper ecosystem integration and staged code execution via a secure cloud computer.

5. What is the scorecard’s most useful measure?

Cost per successful task, including human review, retries and rework.

6. What practical lesson comes from GPT-Red?

Treat external content as untrusted and separate reading from high-impact actions.

7. What do the EU measures change immediately for most users?

Very little today; they set implementation obligations and future competition conditions.

8. What does ASML’s outlook prove?

It confirms the company expects strong chipmaking demand; it does not prove every AI product will succeed.

DO THIS WEEK

Your practical action checklist

- Pick one AI-assisted workflow that happens at least weekly.
- Write a one-sentence definition of “done.”
- Record the approved source data and prohibited data.
- Run five to ten representative examples.
- Measure cost, staff time, retries, correction time and accepted outcomes.
- Add a human approval before any public or irreversible action.
- Test one harmless prompt-injection example if the workflow reads external content.
- Review connector permissions and remove anything unused.
- Document the model/tool version and plan level.
- Create a fallback for service failure.
- Decide: keep, revise, pause or remove the workflow.

Clearforge challenge: do not add a new subscription this week until one existing tool has passed this checklist.

Human accountability

AI can prepare evidence and suggest options. A person remains responsible for accuracy, consent, safety, legal obligations, financial decisions and what is published under the business name.

Seven days of social learning prompts

Monday — Can your AI workflow prove it saves time?

Token price is not the same as task value. This week's guide shows how to count retries, review and rework so you can compare tools by accepted outcomes—not impressive first drafts. What would you include in the real cost of an AI task?

Tuesday — Can Canva really run your creative workflow?

Canva AI 2.0 connects design, research, scheduling and brand rules in a research preview. The opportunity is large, but so is the permission surface. Which part of your creative workflow would you test first?

Wednesday — Would you trust an open-weight model with your workflow?

Inkling brings a very large multimodal model into the customisation conversation. Open weights can increase control, but they also move hosting, evaluation and safety responsibility toward the deployer. Is that extra control worth the operational work for your business?

Thursday — Can AI spot a security problem before code is merged?

GitHub's preview places AI-labelled detections directly inside pull requests. Useful signal, yes—automatic proof of safety, no. Where would an AI warning be most useful in your own approval process?

Friday — Can your research notebook check its own calculations?

Gemini Notebook is gaining a secure cloud computer for code-backed analysis. That can turn sources into tables and charts, but running code does not guarantee the interpretation is right. What calculation would you use as your first controlled test? The weekly Clearforge digest follows the practical results.

Saturday — Could a webpage trick your connected AI?

GPT-Red highlights why prompt injection matters when AI reads browsers, email and files. The practical response is simple: separate research from action and keep approval before anything sensitive. Which action would you never let an AI take without you?

Sunday — Which AI task will you measure next week?

The full Clearforge guide turns this week's releases into one practical checklist. Read it on the Clearforge website, or listen to Clearforge AI Briefing on Spotify and your usual podcast app. Share your questions—or suggest the update you want Clearforge to investigate next.

PRIMARY AND AUTHORITATIVE EVIDENCE

Linked sources

Canva — 16 Jul 2026
Introducing Canva AI 2.0
https://www.canva.com/newsroom/news/canva-create-2026-ai/

Thinking Machines Lab — 15 Jul 2026
Inkling: Our open-weights model
https://thinkingmachines.ai/news/introducing-inkling/

GitHub — 14 Jul 2026
Code scanning shows AI security detections on pull requests
https://github.blog/changelog/2026-07-14-code-scanning-shows-ai-security-detections-on-pull-requests/

OpenAI — 17 Jul 2026
A scorecard for the AI age
https://openai.com/index/a-scorecard-for-the-ai-age/

OpenAI — 15 Jul 2026
GPT-Red: Unlocking Self-Improvement for Robustness
https://openai.com/index/unlocking-self-improvement-gpt-red/

OpenAI — 16 Jul 2026
Why teens deserve access to safe AI
https://openai.com/index/why-teens-deserve-access-safe-ai/

Google — 16 Jul 2026
NotebookLM is now Gemini Notebook
https://blog.google/innovation-and-ai/products/gemini-notebook/notebooklm-gemini-notebook/

European Commission — 16 Jul 2026
AI interoperability on Android and sharing Google Search data under the DMA
https://digital-markets-act.ec.europa.eu/commission-provides-guidance-google-ai-interoperability-android-and-sharing-google-search-data-under-2026-07-16_en

Reuters via Investing.com — 16 Jul 2026
Google required to open up to AI and search rivals under EU-mandated changes
https://www.investing.com/news/stock-market-news/google-required-to-open-up-to-ai-search-engine-rivals-under-eumandated-changes-4795675

ASML — 15 Jul 2026
Q2 2026 financial results and capacity outlook
https://www.asml.com/en/news/press-releases/2026/q2-2026-financial-results

Reuters via Investing.com — 14 Jul 2026
Reflection signs computing deal with Nebius
https://www.investing.com/news/stock-market-news/ai-startup-reflection-signs-over-1-billion-computing-deal-with-nebius-4790506

Source status was rechecked on 20 July 2026. Product access, pricing and implementation dates can change; recheck the primary page before acting.

CLEARFORGE CLOSING

Turn the noise into one useful experiment

The strongest teams will not be the ones that try every release. They will be the ones that choose a real problem, set a quality bar, protect the data, keep a human responsible and measure the result.

This week, choose one workflow and make it observable: what went in, what came out, what it cost, what needed correction and who approved it. That habit will outlast any individual model launch.

Continue learning

Read daily Clearforge briefings, practical features and future weekly reports at clearforge-daily-brief.netlify.app. Listen to the spoken edition through Clearforge AI Briefing on Spotify.

About Clearforge

AI news is noisy. Clearforge turns it into clear, usable daily learning and content. Human-led. AI-empowered. A Sapiver Press production.

Next weekly brief: verify the live sources again before relying on any product, plan or rollout statement.